

Slowpoke: End-to-end Throughput Optimization Modeling for Microservice Applications

Yizheng Xie, Di Jin, Oğuzhan Çölkesen, Vasiliki Kalavri^{BU}, John Liagouris^{BU}, Nikos Vasilakis



github.com/atlas-brown/slowpoke




BROWN


BOSTON
UNIVERSITY

Microservices are Pervasive

Monoliths




Microservices



 **Werner Vogels** 
@Werner

.@Jakewk something like this helps? Real-time graph of microservice dependencies at amazon.com in 2008.




12:48 PM · Jun 11, 2016


Rebuilding Netflix Video Processing Pipeline with Microservices

 Netflix Technology Blog Follow 9 min read · Jan 10, 2024


1.6K 22




Throughput Improvement




Example Shopping Application



Identifying High-Impact Service Optimizations




Identifying High-Impact Service Optimizations




Distributed Tracing Only Captures Current Execution

Call graph for checkout request:



Current execution!



*Slowpoke is a system for **accurately quantifying the effect of a hypothetical optimization** on end-to-end **throughput** for microservice applications.*


Using Slowpoke

Inputs: (1) A target service X.

(2) A hypothetical optimization. E.g., reducing the processing time of 500ms by 50%.

(3) A workload.

Output: End-to-end throughput as if optimization was implemented (Y req/s)



Example:

Input: (1) Cart service

(2) 50% processing time reduction

(3) A mix workload

Output: ~6800 Req/s (from ~3900 Req/s)



Using Slowpoke

Inputs: (1) A target service X.

(2) A hypothetical optimization. E.g., reducing the processing time of 500ms by 50%.

(3) A workload.

Output: End-to-end throughput as if optimization was implemented (Y req/s)




Two Executions and One Performance Model

Inputs: (1) A target service X.



(2) A hypothetical optimization *Eng.*, reducing the processing time of 500ms by 50%.

(3) A workload.

Output: End-to-end throughput as if optimization was implemented (Y req/s)





Applying Causal Profiling to Microservices is Hard



Identifying which threads to slow down incurs **high runtime overhead**.


Slowing down these threads at the right time requires **careful synchronization**.

Key Insight: Slowpoke Slows Down All Non-target Services




Slowpoke slows down ***all non-target services*** with distributed slowdown ***coordination!***


Slowing Down All Non-target Services




Bottleneck Equivalence



(b) Optimized Execution



(c) Slowed-down Execution



Same Bottleneck!

Modeling Bottleneck Equivalence Using Linear Programming

Slowpoke generalizes to:


- Multithreaded services
- Dynamic call graphs
- Different types of bottlenecks, e.g., mutexes.

Maximize end-to-end throughput

Subject to $\left\{ \begin{array}{l} \text{Constraint - A} \\ \text{Constraint - B} \\ \text{Constraint - C} \end{array} \right\}$


E.g., CPU usage $\leq$ CPU quota

(b) Optimized Execution



Linear program (opt)


(c) Slowed-down Execution



Constraint equivalence

Linear program (slowdown)

Ideal Slowdown Scenarios




Pause: signal

Ideal slowdown

High runtime overhead



Naïve Batching Introduces Idle Time



Naïve batching
(E.g., with size=2 req)

Idle time due to serialized slowdown


Slowpoke Uses Distributed Slowdown Coordination



Batching with **coordination**

- Triggered by target service
- Gossip-style propagation
- Paused for batched slowdowns

Data Plane: Instrumented Libraries




Libraries:

- Track batched slowdowns

↔ (1) Data channel: HTTP/gRPC request/response

Data Plane: Poker



- ↔ (1) Data channel: HTTP/gRPC request/response
- ↔ (2) Unix pipe
- ↔ (3) Control channel—TCP
- (4) SIGSTOP/SIGCONT signals

Libraries:

- Track batched slowdowns

Poker:

- Collect batched slowdowns
- Coordinate slowdowns
- Pause services

Evaluation


How accurate is Slowpoke's throughput prediction?

Evaluation Overview

		Benchmarks	Services	LOC
4 applications	{	OnlineBoutique	9	1088
		HotelReservation	6	608
		SocialNetwork	6	532
		MovieReview	12	913
108 synthetic applications ➡		Hydragen ^[1]	3–43	389

$$\text{Error} = \frac{\text{Predicted} - \text{Ground Truth}}{\text{Ground Truth}}$$

9 topologies




3 configurations


- Protocol: e.g., HTTP
- Request concurrency: e.g., sequential
- Target service position: e.g., root in a tree

[1] Mohammad Reza Saleh Sedghpour, Aleksandra Obeso Duque, Xuejun Cai, Björn Skubic, Erik Elmroth, Cristian Klein, and Johan Tordsson. "Hydragen: A microservice benchmark generator." In 2023 IEEE 16th International Conference on Cloud Computing (CLOUD), pp. 189-200. IEEE, 2023.

Prediction Results for Four Applications



Prediction Results for Synthetic Benchmarks






C2—Chain (length 2) F7—Fan-out (width 7) D7—DAG (7 services) M5—Dynamic cycle B*—target is bottleneck
C9—Chain (length 9) D5—DAG (5 services) M4—Dynamic chain M3—Dynamic cache N*—target is not bottleneck
F3—Fan-out (width 3) T8—Unbalanced tree M7—Dynamic tree H*—HTTP, g*—gRPC R/S*—concurrent/sequential


108 Benchmarks × 10 Optimization Percentages

Root mean squared error: 1.89%

More Results in the Paper



- Realistic scaling optimizations
- Mutex-induced bottlenecks
- Large-scale deployment
- A set of microbenchmarks



Slowpoke: End-to-end Throughput Optimization Modeling for Microservice Applications

 github.com/atlas-brown/slowpoke

Yizheng Xie, Di Jin, Oğuzhan Çölkesen, Vasiliki Kalavri^{BU}, John Liagouris^{BU}, Nikos Vasilakis



atlas.cs.brown.edu sites.bu.edu/casp

